


# Examples of Two- Dimensional Systolic Arrays

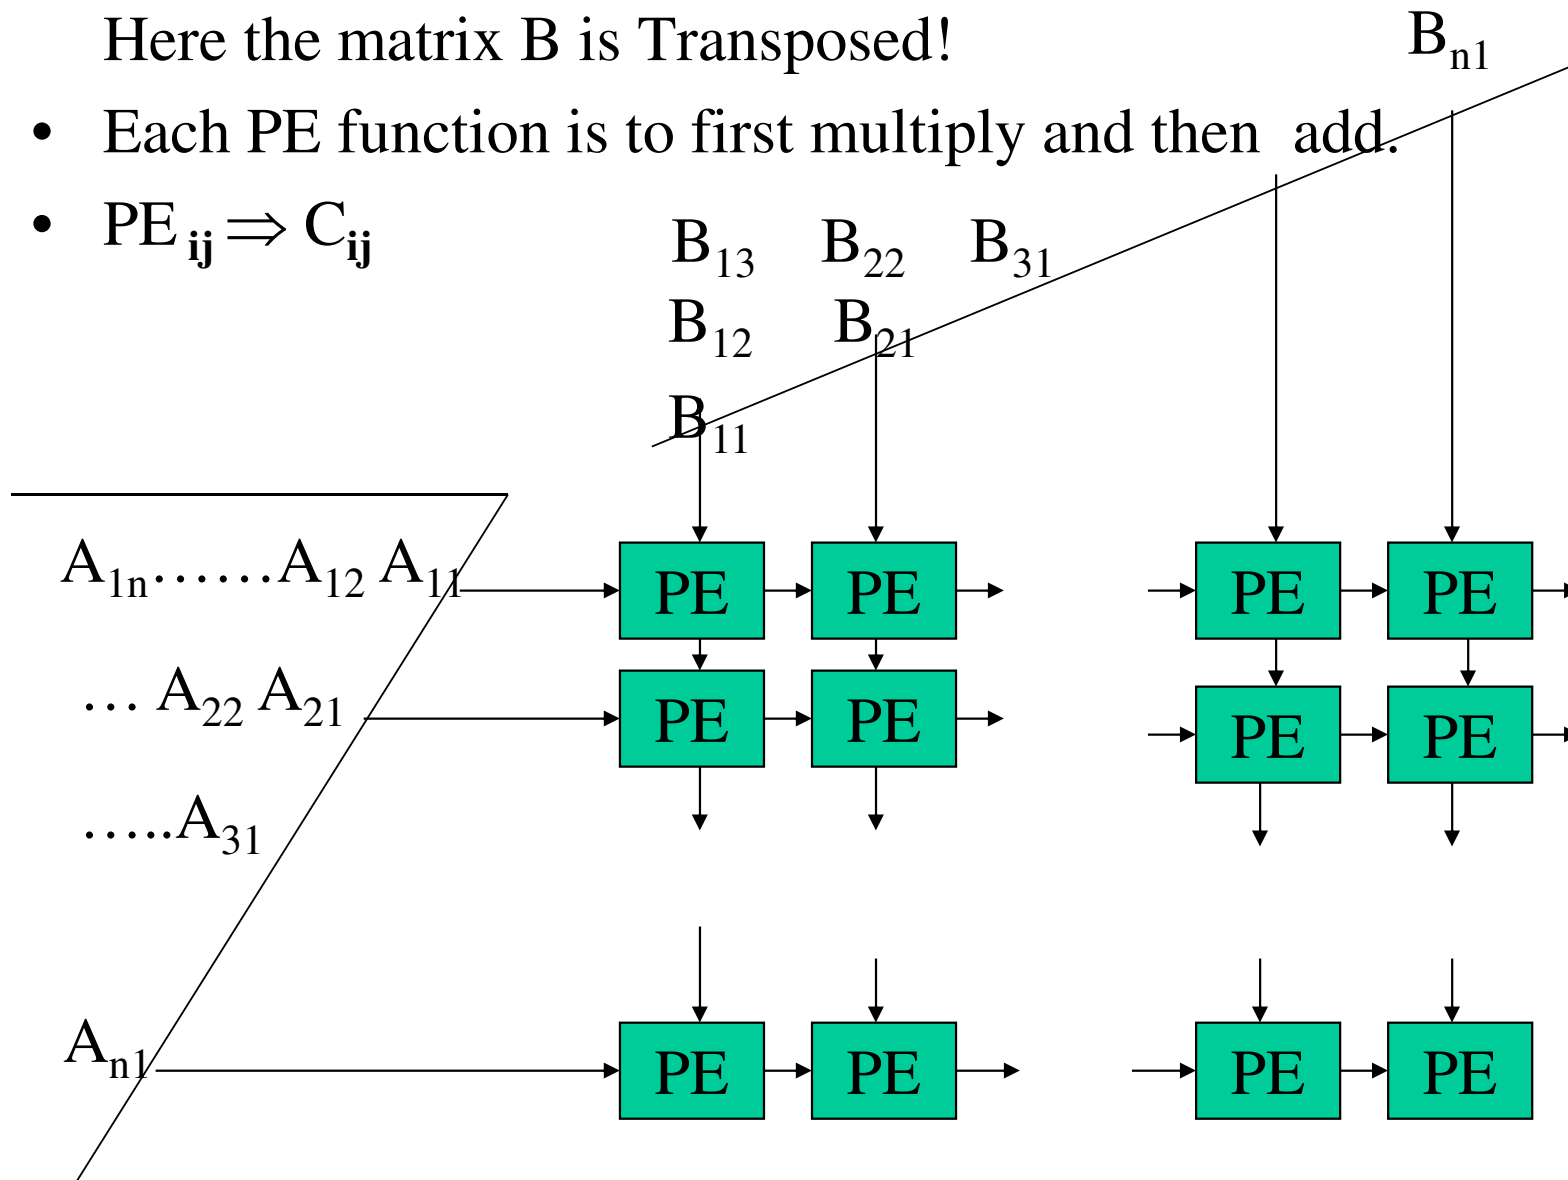
# Obvious Matrix Multiply



- Multiplication

Here the matrix B is Transposed!

- Each PE function is to first multiply and then add.
- $PE_{ij} \Rightarrow C_{ij}$




# Example 4: A Related Algorithm: Cannon's Method

- Let's take another view of systolic multiplication:
  - Consider the rows and columns of the matrices to be multiplied as **strips** that are slide past each other.
- The strips are staggered so that the correct elements are multiplied at each time step.

First step

columns of b (inverted)

rows of a (reversed)




Second step


columns of b (inverted)

rows of a (reversed)

wavefront




## Third step



## Fourth step

columns of b (inverted)

rows of a (reversed)






Fifth step

columns of b (inverted)

rows of a (reversed)




# Cannon's Method

- Rather than have some processors idle,
  - wrap the array rows and columns so that **every processor is doing something on each step.**
- In other words, rather than feeding in the elements, they are rotated around,
  - starting in an initially staggered position as in the systolic model.
- We also change the order of products slightly, to make it correspond to more natural storage by rows and columns.

# Cannon Variation

Note that the a diagonal is in the left column and the b diagonal is in the top row.



Products computed at each step:

$$\text{Example sum: } c_{0,2} = a_{0,2} * b_{2,2} + a_{0,1} * b_{1,2} + a_{0,0} * b_{0,2}$$

Step 1

$a_{00} * b_{00}$	$a_{01} * b_{11}$	$a_{02} * b_{22}$
$a_{11} * b_{10}$	$a_{12} * b_{21}$	$a_{10} * b_{02}$
$a_{22} * b_{20}$	$a_{20} * b_{01}$	$a_{21} * b_{12}$

Step 2

$a_{02} * b_{20}$	$a_{00} * b_{01}$	$a_{01} * b_{12}$
$a_{10} * b_{00}$	$a_{11} * b_{11}$	$a_{12} * b_{22}$
$a_{21} * b_{10}$	$a_{22} * b_{21}$	$a_{20} * b_{02}$

Step 3

$a_{01} * b_{10}$	$a_{02} * b_{21}$	$a_{00} * b_{02}$
$a_{12} * b_{20}$	$a_{10} * b_{01}$	$a_{11} * b_{12}$
$a_{20} * b_{00}$	$a_{21} * b_{11}$	$a_{22} * b_{22}$

# Application of Cannon's Technique

- Consider matrix multiplication of 2  $n \times n$  matrices on a distributed memory machine, on say,  $n^2$  processing elements.
- An obvious way to compute is to think of the PE's as a matrix, with each computing one element of the product.
- We would send each row of the matrix to  $n$  processors and each column to  $n$  processors.
- **In effect, in the obvious way, each matrix is stored a total of  $n$  times.**


# Cannon's Method

- Cannon's method avoids storing each matrix  $n$  times, instead **cycling** (“piping”) the elements through the PE array.
- (It is sometimes called the “pipe-roll” method.)
- The problem is that this cycling is typically too fine-grain to be useful for element-by-element multiply.

# Partitioned Multiplication

- Partitioned multiplication divides the matrices into **blocks**.
- It can be shown that multiplying the individual blocks as if elements of matrices themselves gives the matrix product.

# Block Multiplication



# Cannon's Method is Fine for Block Multiplication

- The blocks are aligned initially as the elements were in our description.
- At each step, entire blocks are transmitted down and to the left of neighboring PE's.
- Memory space is conserved.



# Exercise

- Analyze the running time for the block version of Cannon's method for two  $n \times n$  matrices on  $p$  processors, using  $t_{\text{comp}}$  as the unit operation time and  $t_{\text{comm}}$  as the unit communication time and  $t_{\text{start}}$  as the per-message latency .
- Assume that any pair of processors can communicate in parallel.
- Each block is  $(n/\text{sqrt}(p)) \times (n/\text{sqrt}(p))$ .


## Example 6: Fox's Algorithm

- This algorithm is also for block matrix multiplication; it has a **resemblance to Cannon's algorithm**.
- The difference is that on each cycle:
  - A row block is **broadcast** to every other processor in the row.
  - The column blocks are **rolled cyclically**.

# Fox's Algorithm

Step 1

$a_{00} \cdot b_{00}$	$a_{00} \cdot b_{01}$	$a_{00} \cdot b_{02}$
$a_{11} \cdot b_{10}$	$a_{11} \cdot b_{11}$	$a_{11} \cdot b_{12}$
$a_{22} \cdot b_{20}$	$a_{22} \cdot b_{21}$	$a_{22} \cdot b_{22}$



A different row block of  $a$  is broadcast in each step.

Step 2

$a_{01} \cdot b_{10}$	$a_{01} \cdot b_{11}$	$a_{01} \cdot b_{12}$
$a_{12} \cdot b_{20}$	$a_{12} \cdot b_{21}$	$a_{12} \cdot b_{22}$
$a_{20} \cdot b_{00}$	$a_{20} \cdot b_{01}$	$a_{20} \cdot b_{02}$



$b$  columns are rolled

Step 3

$a_{02} \cdot b_{20}$	$a_{02} \cdot b_{21}$	$a_{02} \cdot b_{22}$
$a_{10} \cdot b_{00}$	$a_{10} \cdot b_{01}$	$a_{10} \cdot b_{02}$
$a_{21} \cdot b_{10}$	$a_{21} \cdot b_{11}$	$a_{21} \cdot b_{12}$

# Synchronous Computations

# Barriers

- Mentioned earlier
- Synchronize all of a group of processes
- Used in both distributed and shared-memory
- Issue: Implementation & cost

# Counter Method for Barriers

- One-phase version
  - Use for distributed-memory
  - Each processor sends a message to the others when barrier reached.
  - When each processor has received a message from all others, the processors pass the barrier

# Counter Method for Barriers

- Two-phase version
  - Use for shared-memory
  - Each processor sends a message to the master process.
  - When the master has received a message from all others, it sends messages to each indicating they can pass the barrier.
  - Easily implemented with blocking receives, or semaphores (one per processor).

# Tree Barrier

- Processors are organized as a tree, with each sending to its parent.
- **Fan-in phase:** When the root of the tree receives messages from both children, the barrier is complete.
- **Fan-out phase:** Messages are then sent down the tree in the reverse direction, and processes pass the barrier upon receipt.



# Butterfly Barrier

- Essentially a fan-in tree for **each** processor, with some sharing toward the leaves.
- Advantage is that no separate fan-out phase is required.

# Butterfly Barrier




Figure 6.6 Butterfly construction.

# Barrier Bonuses

- To implement a barrier, it is only necessary to increment a count (shared memory) or send a couple of messages per process.
- These are communications with null content.
- By adding content to messages, barriers can have added utility.

# Barrier Bonuses

- These can be accomplished along with a barrier:
  - Reduce according to binary operator (esp. good for tree or butterfly barrier)
  - All-to-all broadcast

# Data Parallel Computations

- **forall** statement:

```
forall( j = 0; j < n; j++ )  
    {  
        ... body done in parallel for all j ...  
    }
```

# forall synchronization assumptions

- There are different interpretations of **forall**, so you need to “read the fine print”.
- Possible assumptions from weakest to strongest:
  - No implied synchronization
  - Implied barrier at the end of each loop body
  - Implied barrier before each assignment
  - Each machine instruction synchronized, SIMD-fashion

# Example: Prefix-Sum




Figure 6.8 Data parallel prefix sum operation.

# Example: Prefix-Sum

- Assume that  $n$  is a power of 2.
- Assume shared memory.
- Assume barrier before assignments
- $\text{for}(j = 0; j < \log(n); j++)$

**forall**(  $i = 2^j; i < n; i++$  )

$\underbrace{x[i]}_{\text{effectively buffered new value}} \text{ += } \underbrace{x[i - 2^j]}_{\text{old value}};$

effectively **buffered** new value




# Implementing forall using SPMD:

Assuming PP: “Synchronous Iteration” (barrier at end of body)

- for( j = 0; j < log(n); j++ )  
    forall( i = 0; i < n; i++ )  
        Body(i);

implementable in SPMD as:

- for( j = 0; j < log(n); j++ )  
    {  
    i = my\_process\_rank();  
    Body(i);  
    barrier();  
    }



Outer  
**forall** processes  
implicit

# Example:

## Iterative Linear Equation Solver

```
for( iter = 0; iter < numIterations; iter++ )
```

```
    forall( i = 0; i < n; i++ )
```

```
    {
```

```
        double sum = 0;
```

```
        for( j = 0; j < n; j++ )
```



```
            sum += a[i][j]*x[j];
```

```
        x[i] = sum;
```

```
    }
```

Note:

*Local* memory  
for each i.



# Iterative Linear Equation Solver: Translation to SPMD

```
for( iter = 0; iter < numIterations; iter++ )
```

```
{
```

```
  i = my_process_rank();
```

```
  double sum = 0;
```


```
  for( j = 0; j < n; j++ )
```

```
    sum += a[i][j]*x[j];
```

```
  new_x[i] = sum;
```

```
  all gather new_x to x (implied barrier)
```

```
}
```



Outer  
forall processes  
implicit

# Nested forall's

- ```
for( iter = 0; iter < numIterations; iter++ )  
    forall( i = 0; i < m; i++)  
        forall( j = 0; j < n; i++)  
            Body(i, j)
```

# Example of nested forall's: Laplace Heat equation

- ```
for( iter = 0; iter < numIterations; iter++ )  
  forall( i = 0; i < m; i++)  
    forall( j = 0; j < n; j++)  
      x[i][j] = (x[i-1][j] + x[i][j-1] + x[i+1][j] + x[i][j+1])/4.0;
```

# Exercise

- How would you translate nested forall's to SPMD?

# Synchronous Computations

- Synchronous computations have the form  
    (Barrier)  
    Computation  
    Barrier  
    Computation  
    ...
- Frequency of the barrier and homogeneity of the intervening computations on the processors may vary
- We've seen some synchronous computations already (Jacobi2D, Systolic MM)

# Synchronous Computations

- Synchronous computations can be simulated using asynchronous programming models
  - Iterations can be tagged so that the appropriate data is combined
- Performance of such computations depends on the granularity of the platform, how expensive synchronizations are, and how much time is spent idle waiting for the right data to arrive



# Barrier Synchronizations

- Barrier synchronizations can be implemented in many ways:
  - As part of the algorithm
  - As a part of the communication library
    - PVM and MPI have barrier operations
  - In hardware
- Implementations vary

# Review

- What is time balancing? How do we use time-balancing to decompose Jacobi2D for a cluster?
- Describe the general flow of data and computation in a pipelined algorithm. What are possible bottlenecks?
- What are the three stages of a pipelined program? How long will each take with  $P$  processors and  $N$  data items?
- Would pipelined programs be well supported by SIMD machines? Why or why not?
- What is a systolic program? Would a systolic program be efficiently supported on a general-purpose MPP? Why or why not?

# Common Parallel Programming Paradigms


- Embarrassingly parallel programs
- Workqueue
- Master/Slave programs
- Monte Carlo methods
- Regular, Iterative (Stencil) Computations
- Pipelined Computations
- Synchronous Computations

# Synchronous Computations

- Synchronous computations are programs structured as a group of separate computations which must at times wait for each other before proceeding
- Fully synchronous computations = programs in which all processes synchronized at regular points
- Computation between synchronizations often called stages

# Synchronous Computation Example: Bitonic Sort

- Bitonic Sort an interesting example of a synchronous algorithm
- Computation proceeds in stages where each stage is a (smaller or larger) **shuffle-exchange network**
- **Barrier synchronization** at each stage



# Bitonic Sort

- A **bitonic sequence** is a list of keys

$a_0, a_1, \dots, a_{n-1}$  such that


- 1 For some  $i$ , the keys have the ordering

$$a_0 \leq a_1 \leq \dots \leq a_i \geq \dots \geq a_{n-1}$$

or

- 2 The sequence can be shifted cyclically so that

1) holds



# Bitonic Sort Algorithm

- The bitonic sort algorithm recursively calls two procedures:
  - **BSORT**(i,j,X) takes **bitonic** sequence  $a_i, a_{i+1}, \dots, a_j$  and produces a non-decreasing (X=+) or a non-increasing **sorted** sequence (X=-)
  - **BITONIC**(i,j) takes an **unsorted** sequence  $a_i, a_{i+1}, \dots, a_j$  and produces a **bitonic** sequence
- The main algorithm is then
  - **BITONIC**(0,n-1)
  - **BSORT**(0,n-1,+)

# How does it do this?

- We'll show how BSORT and BITONIC work but first consider an interesting property of bitonic sequences:

Assume that  $a_0, a_1, \dots, a_{n-1}$  is bitonic and that  $n$  is even. Let

$$B_1 = \min\{a_0, a_{\frac{n}{2}}\}, \min\{a_1, a_{\frac{n}{2}+1}\}, \dots, \min\{a_{\frac{n}{2}-1}, a_{n-1}\}$$

$$B_2 = \max\{a_0, a_{\frac{n}{2}}\}, \max\{a_1, a_{\frac{n}{2}+1}\}, \dots, \max\{a_{\frac{n}{2}-1}, a_{n-1}\}$$

Then  $B_1$  and  $B_2$  are bitonic sequences and for all

$$x \in B_1, y \in B_2 \quad x \leq y$$





# Picture “Proof” of Interesting Property

- Consider

$$B_1 = \min\{a_0, a_{\frac{n}{2}}\}, \min\{a_1, a_{\frac{n}{2}+1}\}, \dots, \min\{a_{\frac{n}{2}-1}, a_{n-1}\}$$

$$B_2 = \max\{a_0, a_{\frac{n}{2}}\}, \max\{a_1, a_{\frac{n}{2}+1}\}, \dots, \max\{a_{\frac{n}{2}-1}, a_{n-1}\}$$

- Two cases:  $i \geq \frac{n}{2}$  and  $i < \frac{n}{2}$





# Picture “Proof” of Interesting Property

- Consider

$$B_1 = \min\{a_0, a_n\}, \min\{a_1, a_{\frac{n}{2}+1}\}, \dots, \min\{a_{\frac{n}{2}-1}, a_{n-1}\}$$

$$B_2 = \max\{a_0, a_n\}, \max\{a_1, a_{\frac{n}{2}+1}\}, \dots, \max\{a_{\frac{n}{2}-1}, a_{n-1}\}$$

$$i \geq \frac{n}{2}$$



$$x \in B_1, y \in B_2$$

$$x \leq y$$


# Picture “Proof” of Interesting Property

- Consider

$$B_1 = \min\{a_0, a_{\frac{n}{2}}, \min\{a_1, a_{\frac{n}{2}+1}\}, \dots, \min\{a_{\frac{n}{2}-1}, a_{n-1}\}\}$$


$$B_2 = \max\{a_0, a_{\frac{n}{2}}, \max\{a_1, a_{\frac{n}{2}+1}\}, \dots, \max\{a_{\frac{n}{2}-1}, a_{n-1}\}\}$$

$$i < \frac{n}{2}$$



# Back to Bitonic Sort


- Remember
  - **BSORT**(i,j,X) takes bitonic sequence  $a_i, a_{i+1}, \dots, a_j$  and produces a non-decreasing (X=+) or a non-increasing sorted sequence (X=-)
  - **BITONIC**(i,j) takes an unsorted sequence  $a_i, a_{i+1}, \dots, a_j$  and produces a bitonic sequence
- Let's look at BSORT first ...



# Here's where the shuffle-exchange comes in ...

- Shuffle-exchange network routes the data correctly for comparison
- At each shuffle stage, can use + switch to separate B1 and B2


$$\begin{array}{c} a \\ b \end{array} \begin{array}{|c|} \hline + \\ \hline \end{array} \begin{array}{l} \min\{a,b\} \\ \max\{a,b\} \end{array}$$



# Sort bitonic subsequences to get a sorted sequence

- **BSORT( $i, j, X$ )**

- If  $|j-i| < 2$  then return  $[\min(i, i+1), \max(i, i+1)]$
- Else
  - Shuffle( $i, j, X$ )
  - Unshuffle( $i, j$ )
  - Pardo
    - BSORT( $i, i+(j-i+1)/2 - 1, X$ )
    - BSORT( $i+(j-i+1)/2 + 1, j, X$ )



# BITONIC takes an unsorted sequence as input and returns a bitonic sequence

- **BITONIC(i,j)**

- If  $|j-i| < 2$  then return  $[i, i+1]$



- Else

- Pardo

- BITONIC( $i, i+(j-i+1)/2 - 1$ ); BSORT ( $i, i+(j-i+1)/2 - 1, +$ )


- BITONIC( $i+(j-i+1)/2 + 1, j$ ); BSORT ( $i+(j-i+1)/2 + 1, j, -$ )

(note that any 2 keys are already a bitonic sequence)



# Putting it all together

- Bitonic sort for 8 keys:





# Complexity of Bitonic Sort

$$\begin{aligned}T_{BSORT}(n = 2^j) &= 2 + T(2^{j-1}) \\&= \dots = 2(j-1) + T(2) = O(j) \\&= O(\log n)\end{aligned}$$

$$\begin{aligned}T_{BITONIC}(n = 2^j) &= T(2^{j-1}) + 2(j-1) + 1 \\&= \sum_{i=2}^{j-1} 2(i-1) + 1 \\&= O(j^2) \\&= O(\log^2 n)\end{aligned}$$

# Programming Issues

- Flow of data is assumed to transfer from stage to stage synchronously; usual issues with performance if algorithm is executed asynchronously
- Note that logical interconnect for each problem size is different
  - Bitonic sort must be mapped efficiently to target platform
- Unless granularity of platform very fine, multiple comparators will be mapped to each processor

# Review


- What is a a synchronous computation?  
What is a fully synchronous computation?
- What is a bitonic sequence?
- What do the procedures BSORT and BITONIC do?
- How would you implement Bitonic Sort in a performance-efficient way?

# Mapping Bitonic Sort

For every stage the 2X2 switches compare keys which differ in a single bit

$$a_k \cdots a_{i+1} 1 a_{i-1} \cdots a_0$$

$$a_k \cdots a_{i+1} 0 a_{i-1} \cdots a_0$$





# Supporting Bitonic Sort on a hypercube

- Switch comparisons can be performed in constant time on hypercube interconnect.


$$a_k \cdots a_{i+1} 1 a_{i-1} \cdots a_0$$

$$a_k \cdots a_{i+1} 0 a_{i-1} \cdots a_0$$



# Mappings of Bitonic Sort

- Bitonic sort on a multistage full shuffle
  - Small shuffles do not map 1-1 to larger shuffles!
  - Stone used a clever approach to map logical stages into full-sized shuffle stages while preserving  $O(\log^2 n)$  complexity




# Outline of Stone's Method

- Pivot bit = index being shuffled
- Stone noticed that for successive stages, the pivot bits are  $\underline{a_0}, \underline{a_1}, \underline{a_0}, \underline{a_2}, \underline{a_1}, \underline{a_0}, \dots, \underline{a_i}, \underline{a_{i-1}}, \dots, \underline{a_1}, \underline{a_0}, \dots$
- If the pivot bit is in place, each subsequent stage can be done using a full-sized shuffle (a\_0 done with a single comparator)
- For pivot bit j, need k-j full shuffles to position bit j for comparison  $a_k \cdots a_{j+1} a_j \cdots a_0 \rightarrow a_j \cdots a_0 a_k \cdots a_{j+1}$
- Complexity of Stone's method:

$$\begin{aligned} & k + [(k-1) + 1] + [(k-2) + 1 + 1] + \cdots + [(k-k) + 1 + \cdots + 1] \\ & = k(k+1) = O(k^2) = O(\log^2 n) \end{aligned}$$

# Many-one Mappings of Bitonic Sort

- For platforms where granularity is coarser, it will be more cost-efficient to map multiple comparators to one processor
- Several possible conventional mappings




- Compare-split provides another approach ...





# Compare-Split

- For a block of keys, may want to use a **compare-split** operation (rather than compare-exchange) to accommodate multiple keys at a processor
- Idea is to assume that each processor is assigned a **block of keys**, rather than 2 keys
  - **Blocks are already sorted with a sequential sort**
  - To perform compare-split, a processor compares blocks and returns the smaller half of the aggregate keys as the **min block** and the larger half of the aggregate keys as the **max block**



# Compare-Split

- Each Block represents more than one datum
  - Blocks are already sorted with a sequential sort
  - To perform compare-split, a processor compares blocks and returns the smaller half of the aggregate keys as the **min block** and the larger half of the aggregate keys as the **max block**



# Performance Issues

- What is the complexity of compare-split?
- How do we optimize compare-split?
  - How many datum per block?
  - How to allocate blocks per processors?
  - How to synchronize intra-block sorting with inter-block communication?

# Conclusion on Systolic Arrays

- Advantages of systolic arrays are:
  - 1. Regularity and modular design(Perfect for VLSI implementation).
  - 2. Local interconnections(Implements algorithms locality).
  - 3. High degree of pipelining.
  - 4. Highly synchronized multiprocessing.
  - 5. Simple I/O subsystem.
  - 6. Very efficient implementation for great variety of algorithms.
  - 7. High speed and Low cost.
  - 8. Elimination of global broadcasting and modular expansibility.

# Disadvantages of systolic arrays

- The main disadvantages of systolic arrays are:
  - 1. Global synchronization limits due to signal delays.
  - 2. High bandwidth requirements both for periphery(RAM) and between PEs.
  - 3. Poor run-time fault tolerance due to lack of interconnection protocol.

# Parallel overhead.

Food for  
thought



- Running time for a program running on several processors including an allowance for parallel overhead compared with the ideal running time.
- There is often a point beyond which adding further processors doesn't result in further efficiency.
- There may also be a point beyond which adding further processors results in slower execution.

# Sources

- 1. Seth Copen Goldstein, CMU**
- 2. David E. Culler, UC. Berkeley,**
- 3. Keller@cs.hmc.edu**
- 4. Syeda Mohsina Afroze**  
**and other students of Advanced Logic Synthesis,**  
**ECE 572, 1999 and 2000.**
- 5. Berman**